

Final Review

Exam logistics

- Friday, August 17th
- 7-10pm
- Skilling Auditorium
- SCPD students may take the exam any time Thursday or Friday
- Email me if you need any accommodations

Study strategies

- Read through lecture notes
 - Things should Make Sense™
- Revisit assignments
 - Know why everything works the way it does. I'll almost certainly reference your assignments or ask you to extend them
 - May be helpful to put core pieces of code on your cheat sheet
- Work through practice exams
- Work through lab questions

Material

- Filesystems
- Multiprocessing
 - Signals
 - Virtual memory
 - Scheduling
- Multithreading
 - Thread management
 - Synchronization
 - Design decisions
- Networking
 - IP addresses and port numbers
 - DNS
 - Understanding sockets and network connections
 - socket, bind, listen, connect syscalls
 - HTTP protocol
- Other topics
 - Nonblocking I/O
 - MapReduce
 - Guest talks

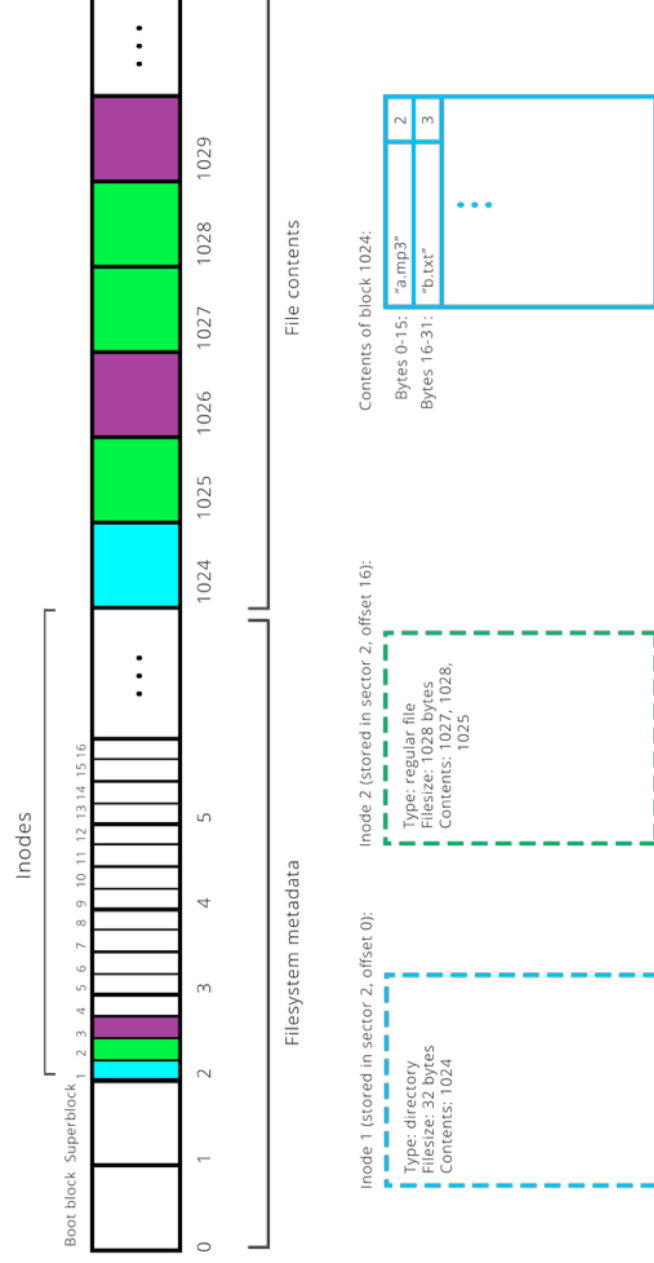
Material

- Filesystems
- Multiprocessing
 - Signals
 - Virtual memory
 - Scheduling
- **Multithreading**
 - Thread management
 - Synchronization
 - Design decisions
- **Networking**
 - IP addresses and port numbers
 - DNS
 - Understanding sockets and network connections
 - socket, bind, listen, connect syscalls
 - HTTP protocol
- Other topics
 - Nonblocking I/O
 - MapReduce
 - Guest talks

Filesystems

- FS design/layering: inodes, files, directories, links
- FS usage: file descriptors, file entries, nodes

Filesystem layout



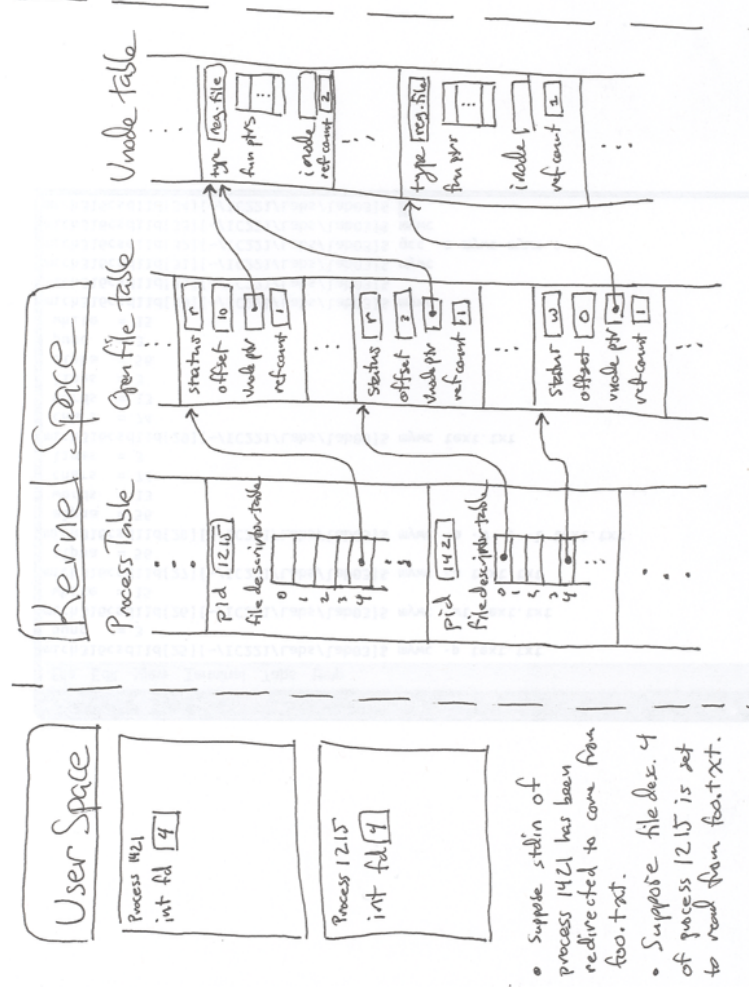
Directories

- Directories are just a special type of file
- Payload consists of (inode, name) pairs
- Know how the path resolution process

Links

- Hard links are directory entries pointing to a file
- Soft links are special files whose payload is a path to a file

File descriptor/file entry/vnode tables



File descriptor/file entry/vnode tables

- What do dup, dup2, open, close do to the tables?
- What is/isn't shared across processes?
- How do pipes work?

More multiprocessing

Signals, handlers, waitpid, sigsuspend
Scheduling

Signals

- A form of inter-process communication
- SIGINT, SIGTSTP, SIGCONT, SIGSTOP, SIGKILL
- `signal(SIGCHLD, reapChild)`
- Signal handlers are per-process and exist in the code segment (they are preserved across `fork` but not across `execvp`)
- Signals are handled when the process is on the CPU
 - If a process is in the blocked set, it will be moved to the ready queue upon receipt of a signal, then (usually) moved back to the blocked set when the signal is handled
- If a SIGCHLD arrives while executing the SIGCHLD handler, delivery of the second signal will be deferred until the handler finishes handling the first signal

Signal-related functions

- `kill`, `raise`, `sigprocmask`, `signal`, `sigsuspend`
- You can use `sigprocmask` to defer signals
- Signal masks are preserved across `fork` and `execvp`
- `Sigsuspend` sleeps until a signal is delivered
 - `sigsuspend(&mask)` :
 //ATOMICALLY:
 `sigset_t old;`
 `sigprocmask(SIG_SETMASK, &mask, &old);`
 `sleep();` // wait for signal to wake us up
 `sigprocmask(SIG_SETMASK, &old, NULL)`

Signal puzzle

```
static pid_t pid;
static int counter = 0;

static void parentHandler(int unused) {
    counter += 2;
    printf("counter = %d\n", counter);
}

static void childHandler(int unused) {
    counter += 1;
    printf("counter = %d\n", counter);
    kill(getppid(), SIGUSR1);
}

1. Can this program DEADLOCK?
BONUS: How many outputs are there? }
```

```
int main(int argc, char *argv[]) {
    signal(SIGUSR1, parentHandler);
    if ((pid = fork()) == 0) {
        signal(SIGUSR1, childHandler);
        sigset_t mask; sigemptyset(&mask);
        sigsuspend(&mask);
        return 0;
    }

    sleep(1); // hmmm...
    kill(pid, SIGUSR1);
    waitpid(pid, NULL, 0);
    counter += 3;
    printf("counter = %d\n", counter);
    return 0;
}
```

Signal puzzle

```
static pid_t pid;
static int counter = 0;

static void parentHandler(int unused) {
    counter += 2;
    printf("counter = %d\n", counter);
}

static void childHandler(int unused) {
    counter += 1;
    printf("counter = %d\n", counter);
    kill(getppid(), SIGUSR1);
}

1. Can this program DEADLOCK?
BONUS: How many outputs are there? }
```

```
int main(int argc, char *argv[]) {
    signal(SIGUSR1, parentHandler);
    if ((pid = fork()) == 0) {
        signal(SIGUSR1, childHandler);
        sigset_t mask; sigemptyset(&mask);
        sigsuspend(&mask);
        return 0;
    }

    sleep(1); // hmmm...
    kill(pid, SIGUSR1);
    waitpid(pid, NULL, 0);
    counter += 3;
    printf("counter = %d\n", counter);
    return 0;
}
```

3 Outputs:

- Output 1. - The child prints 2, and both the child and parent deadlock.
- Output 2. - The child prints 1, the parent prints 2, and both the child and parent deadlock.
- Output 3. - The child prints 1, the parent prints 2, and then the parent prints 5. Both processes exit.

Scheduling

- Process control block: struct representing a process's state
 - PID, register values, file descriptor table, performance statistics, etc
- Running set, ready queue, blocked set
- What causes a process to move from one queue to another?

Multithreading

Threads vs processes
Synchronization: locks, semaphores, condition variables
Design decisions

Threads vs processes

- Threads
 - Lightweight
 - Easier to synchronize and share information
 - Easier to make mistakes
- Processes
 - OS provides isolation and security
 - Harder to communicate and synchronize

Locks and lock guards

- Mutex motivation
 - Prevent race conditions: secure access to shared data structures
- Mutex gotchas:
 - Program can deadlock if you forget to unlock
 - Program can deadlock if you have too many locks and have circular dependencies
 - Program may run slower than necessary if you have too few locks or hold them for too long

Locks and lock guards

- `mutex m;`
 - Constructs mutex in unlocked state
- `m.lock ();`
 - If unlocked, secures lock and proceeds
 - If locked, does not get lock and blocks
- `m.unlock ();`
 - Should only call if you have the lock :D
 - Everyone else waiting on this lock wakes up and tries to acquire it
- `lock_guard <mutex> lg (m);`
 - Same as `m.lock`, except will automatically unlock when it goes out of scope

Semaphores

- Semaphore motivation:
 - “Bucket of balls” analogy
 - Easy primitive to sleep when we need to wait for something, and wake up when it becomes available
- Semaphore gotchas:
 - Program can deadlock if you take something from the bucket and forget to put it back

Semaphores

- semaphore `s` (initial `val`);
 - Constructs semaphore with “initial `val`” balls in the bucket
 - This is *not* a maximum size of the bucket
- `s.wait()`;
 - If `val > 0`, atomically decrements `val` and proceeds
 - If `val == 0`, blocks
- `s.signal()`;
 - “Returns” ball to the bucket by atomically incrementing `val`
 - Potentially wakes up threads which have blocked on `s.wait()` so they can try again :)

Condition variables

- Motivation
 - Wait for some condition to become true
 - More flexible than a semaphore

Condition variables

```
while (!predicate) {  
    wait();  
}
```

Condition variables

```
m.lock();  
  
while (!predicate) {  
    m.unlock();  
    wait();  
    m.lock();  
}  
  
m.unlock();
```

Condition variables

```
cv.wait(m, predicate)
{
    m.lock();
    while (!predicate) {
        // ATOMICALLY:
        m.unlock();
        wait();
        m.lock();
        // END ATOMICALLY
    }
    m.unlock();
}
```

Condition variables

- Condition_variable_any cv;
- Condition variable constructor
- cv.wait (m, predicate);
 - Uses mutex m to safely evaluate whether predicate is true or false
 - If false, blocks until woken up
- cv.notify_one (); cv.notify_all ();
 - Wake up one, or all (depends on which one you call) threads blocked bc of cv.wait. This only wakes them up so they can re-evaluate the predicate--if the predicate is false, they'll go back to sleep!

Condition variables

- Motivation
 - Wait for some condition to become true
 - More flexible than a semaphore
- Gotchas
 - You need to pass a *single* lock that protects any variables in the predicate


Design decisions

- How many threads should you spawn?
 - It depends. CPU-heavy or not?

Networking

- IP addresses and port numbers
- DNS: how resolution works, `gethostbyname()`
- Understanding sockets and network connections
 - socket, bind, listen, connect syscalls
 - HTTP protocol

DNS resolution



Sockets

- **Communicating between processes on a machine:**
 - Use pipes
- **Communicating between different machines:**
 - We'd like to keep using the same kinds of abstractions
 - Use sockets!

Sockets

- Socket descriptors are returned by the `socket()` and `accept()` syscalls
- Nearly interchangeable with file descriptors
- Bidirectional
- Can be used to talk between processes on different machines
- Can be used to establish interprocess communication even after a process has started running

Working with sockets

- Since they act like file descriptors, we can use the read/write/close syscalls
- In practice, we more often use the `socket` and `iosocketstream` abstractions

time-client

```
int main(int argc, char* argv[]) {  
    int client = createClientSocket("myth51.stanford.edu",  
                                   12345);  
    socket sb(client);  
    iosockstream ss(&sb);  
    string timestr;  
    getline(ss, timestr);  
    cout << timestr << endl;  
    return 0;  
}
```

time-client

```
int createClientSocket (const string& host, unsigned short port)
{
```

```
}
```

time-client

```
int createClientSocket (const string& host, unsigned short port)
{
    struct hostent *he = gethostbyname(host.c_str());
```

```
}
```

time-client

```
int createClientSocket (const string& host, unsigned short port)
{
    struct hostent *he = gethostbyname(host.c_str());
    int client = socket(AF_INET, SOCK_STREAM, 0);
```

```
}
```

time-client

```
int createClientSocket (const string& host, unsigned short port)
{
    struct hostent *he = gethostbyname(host.c_str());
    int client = socket(AF_INET, SOCK_STREAM, 0);

    struct sockaddr_in serverAddress;
    memset(&serverAddress, 0, sizeof(serverAddress));
    serverAddress.sin_family = AF_INET;
    serverAddress.sin_port = htons(port);
    serverAddress.sin_addr.s_addr = (struct in_addr *)he->h_addr->s_addr;
```

```
}
```

time-client

```
int createClientSocket (const string& host, unsigned short port)
{
    struct hostent *he = gethostbyname(host.c_str());
    int client = socket(AF_INET, SOCK_STREAM, 0);

    struct sockaddr_in serverAddress;
    memset(&serverAddress, 0, sizeof(serverAddress));
    serverAddress.sin_family = AF_INET;
    serverAddress.sin_port = htons(port);
    serverAddress.sin_addr.s_addr = (struct in_addr *)he->h_addr->s_addr;

    connect(client, (struct sockaddr *) &serverAddress, sizeof(serverAddress));
}
}
```

time-client

```
int createClientSocket (const string& host, unsigned short port)
{
    struct hostent *he = gethostbyname(host.c_str());
    int client = socket(AF_INET, SOCK_STREAM, 0);

    struct sockaddr_in serverAddress;
    memset(&serverAddress, 0, sizeof(serverAddress));
    serverAddress.sin_family = AF_INET;
    serverAddress.sin_port = htons(port);
    serverAddress.sin_addr.s_addr = (struct in_addr *)he->h_addr->s_addr;

    connect(client, (struct sockaddr *) &serverAddress, sizeof(serverAddress));

    return client;
}
```

time-server

```
int main(int argc, char* argv[]) {  
    int server = createServerSocket(12345);  
    ThreadPool pool(8);  
    while (true) {  
        int client = accept(server, NULL, NULL);  
        pool.schedule([client] { publish(client); });  
    }  
    return 0;  
}
```

time-server

```
int createServerSocket(unsigned short port) {
```

```
}
```

time-server

```
int createServerSocket(unsigned short port) {  
    int serverSocket = socket(AF_INET, SOCK_STREAM, 0);
```

```
}
```

time-server

```
int createServerSocket(unsigned short port) {  
    int serverSocket = socket(AF_INET, SOCK_STREAM, 0);  
  
    struct sockaddr_in serverAddress;  
    memset(&serverAddress, 0, sizeof(serverAddress));  
    serverAddress.sin_family = AF_INET;  
    serverAddress.sin_addr.s_addr = htonl(INADDR_ANY);  
    serverAddress.sin_port = htons(port);
```

```
}
```

time-server

```
int createServerSocket(unsigned short port) {
    int serverSocket = socket(AF_INET, SOCK_STREAM, 0);

    struct sockaddr_in serverAddress;
    memset(&serverAddress, 0, sizeof(serverAddress));
    serverAddress.sin_family = AF_INET;
    serverAddress.sin_addr.s_addr = htonl(INADDR_ANY);
    serverAddress.sin_port = htons(port);

    bind(serverSocket, (struct sockaddr *) &serverAddress,
        sizeof(serverAddress));

}
```

time-server

```
int createServerSocket(unsigned short port) {
    int serverSocket = socket(AF_INET, SOCK_STREAM, 0);

    struct sockaddr_in serverAddress;
    memset(&serverAddress, 0, sizeof(serverAddress));
    serverAddress.sin_family = AF_INET;
    serverAddress.sin_addr.s_addr = htonl(INADDR_ANY);
    serverAddress.sin_port = htons(port);

    bind(serverSocket, (struct sockaddr *) &serverAddress,
        sizeof(serverAddress));
    listen(serverSocket, 128);

}
```


time-server

```
int createServerSocket(unsigned short port) {
    int serverSocket = socket(AF_INET, SOCK_STREAM, 0);

    struct sockaddr_in serverAddress;
    memset(&serverAddress, 0, sizeof(serverAddress));
    serverAddress.sin_family = AF_INET;
    serverAddress.sin_addr.s_addr = htonl(INADDR_ANY);
    serverAddress.sin_port = htons(port);

    bind(serverSocket, (struct sockaddr *) &serverAddress,
        sizeof(serverAddress));
    listen(serverSocket, 128);
    return serverSocket;
}
```

Socket API calls

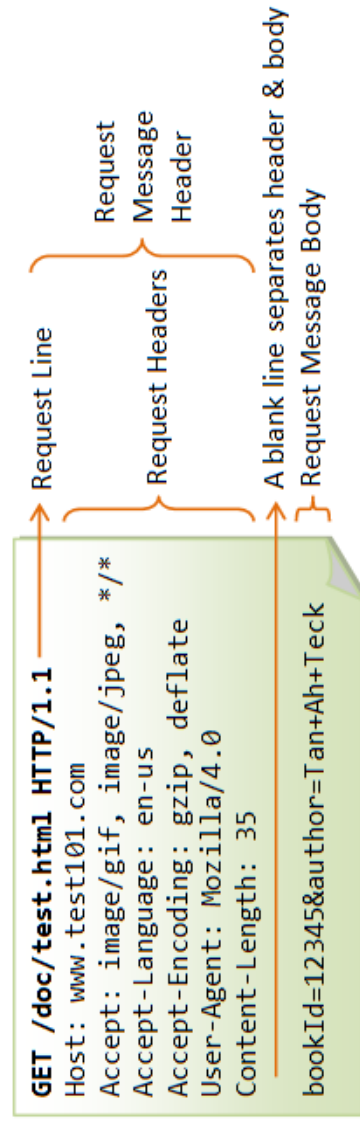
Used by both client and server	Used by server
Socket	Bind
<ul style="list-style-type: none">• Create an endpoint for communication• Returns file descriptor that you can use to create sockbuf, iostream	<ul style="list-style-type: none">• Bind a name to a socket
	Listen
	<ul style="list-style-type: none">• Listen for connections on a socket
	Accept
	<ul style="list-style-type: none">• Waits until someone “rings up” the server• Returns the fd of the client who put in a request
Used by client	
Connect	
<ul style="list-style-type: none">• Initiate a connection on a socket	

Networking questions

- What are the similarities and differences between sockets and pipes?
- Why do we need the reentrant `gethostbyname_r`?
- Which socket API calls could block?
- Why do we handle requests in separate threads?
- How many open, yet-to-be-accepted requests can one server maintain? What about open connections (after accepting)?

HTTP requests/responses

Just know what they look like



Nonblocking I/O

Understand at a conceptual level
What's the point?
What do the epoll functions do?
Edge triggering vs level triggering

Nonblocking I/O

- We've overcome latency on blocking I/O operations by using threads
- However...
 - Threads are expensive
 - We only have a limited number of them
- Alternative: Nonblocking I/O
 - Configure file/socket descriptors as nonblocking descriptors. If we call read() or write() on them, those sys calls will return immediately instead of blocking

Nonblocking I/O

- Writing a nonblocking server:
 - Every time we accept() an incoming network request, configure the descriptor to be nonblocking
 - Add the descriptor to a vector
 - Loop over the vector, calling read() on each file descriptor to check whether there is any new information that has come in on each file descriptor
 - read() returns -1 with errno EAGAIN if there's no data to read right now

Nonblocking I/O

- Writing a nonblocking server:
 - Every time we accept() an incoming network request, configure the descriptor to be nonblocking
 - Add the descriptor to a vector
 - Loop over the vector, calling read() on each file descriptor to check whether there is any new information that has come in on each file descriptor
- Isn't this busy waiting? (Yes.)
 - Solve this problem with the epoll library
 - Similar to waitpid in some sense

Nonblocking I/O

- `epoll_create`: creates a “watch set” of file descriptors
 - Similar to how `sigset_t` is a set of signals you’re waiting for
- `epoll_ctl`: modifies a watch set
 - Similar to `sigprocmask`
- `epoll_wait`: waits until there is activity on a file descriptor in the watch set
 - Similar to `sigsuspend`

Nonblocking I/O questions

- Why does HTTP long polling not work well with your proxy implementation?
- How could you make it better?
- What are some disadvantages of nonblocking I/O?
- Would nonblocking I/O be a good addition for `InternetArchive`?

MapReduce

- Know your implementation
- Make sure you understand the point

Principles of System Design

- Know what each principle means (but don't memorize them)
- Be able to give examples of each principle

Guest talks

Have listened to them
(I won't ask anything obscure)

Take a deep breath! It's going to be okay.

Don't forget to sleep!